

# Profiling and Performance Tuning

## Introduction

This lab guides you through the process of profiling an application and analyzing the output. The application is then accelerated in hardware and profiled again to analyze the performance improvement.

## Objectives

After completing this lab, you will be able to:

- Setup the board support package (BSP) for profiling an application
- Set the necessary compiler directive on an application to enable profiling
- Setup the profiling parameters
- Profile an application and analyze the output

## Procedure

This lab is separated into steps that consist of general overview statements that provide information on the detailed instructions that follow. Follow these detailed instructions to progress through the lab.

## Design Description

In this lab, you will design an embedded system consists of ARM Cortex-A9 processor SoC and two instances of the provided FIR filter IP. The following diagram represents the completed design (**Figure 1**).

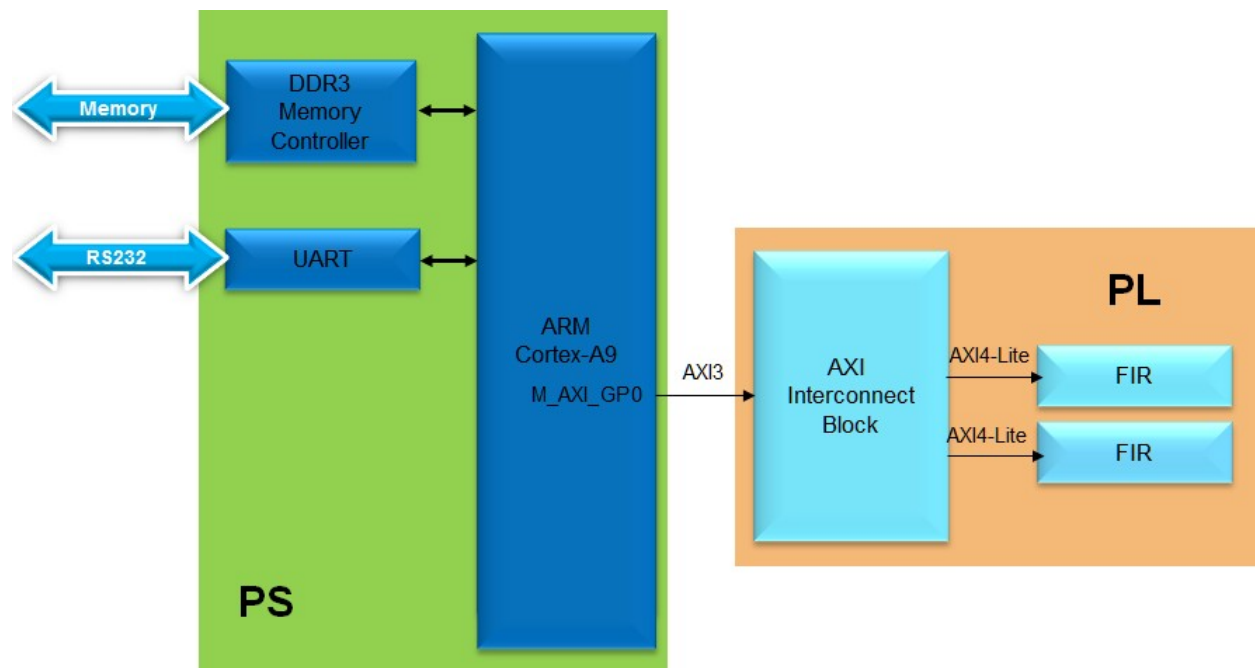
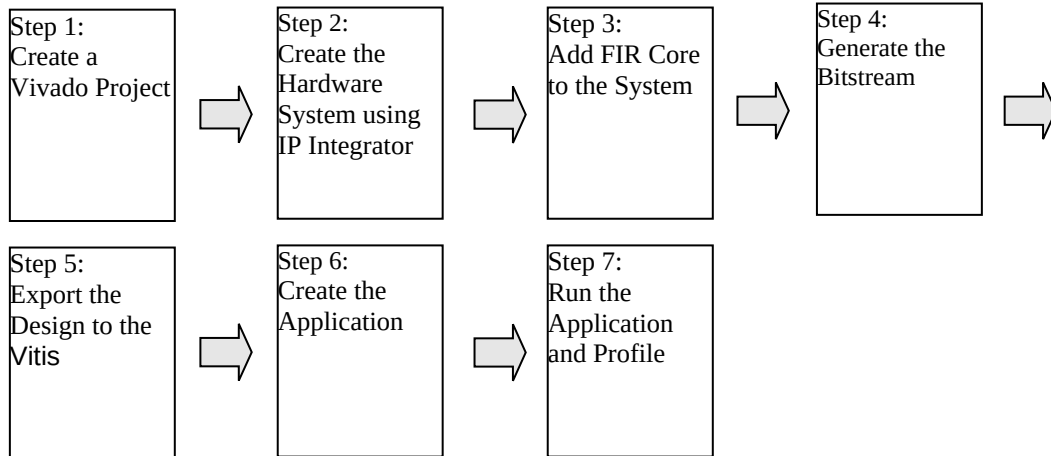


Figure 1. Completed Design

## General Flow for this Lab



### 1 Create a Vivado Project

### Step 1

#### 1.1 Launch Vivado and create an empty project, called lab, targeting the Zybo or ZedBoard Zynq Evaluation and Development Kit and using the VHDL language.

- 1.1.1 Open Vivado and create a new project new project call *lab* in your hardware directory.
- 1.1.2 Select the **RTL Project** option in the *Project Type* form, and click **Next**.
- 1.1.3 Select **VHDL** as the *Target Language* in the *Add Sources* form, and click **Next**.
- 1.1.4 Click **Next** two times.
- 1.1.5 In the *Default Part* form, click on *Boards* and select either the Zybo or Zedboard and click **Next**.
- 1.1.6 Click **Finish** to create an empty Vivado project.

#### 1.2 Set the project settings to include provided fir\_top IP

- 1.2.1 Click **Project Settings** in the *Flow Navigator* pane.
- 1.2.2 Select **IP** in the left pane of the *Project Settings* form.
- 1.2.3 Click on the **Green Plus** button, browse to **fir\_ip** and click **Select**. **fir\_ip** is provided as a source for this lab.
- 1.2.4 The *fir\_top\_v1\_0* IP will appear the **IP in the Selected Repository** window.
- 1.2.5 Click **OK**.

## 2Creating the Hardware System Using IP Integrator

## Step 2

### 2.1 Create a block design in the Vivado project using IP Integrator to generate the Zynq ARM Cortex-A9 processor based hardware system.

2.1.1 In the Flow Navigator, click **Create Block Design** under IP Integrator.

2.1.2 Name the block **system** and click **OK**.

2.1.3 Click on Add IP in the message at the top of the *Diagram* panel.

2.1.4 Once the IP Catalog is open, find **ZYNQ7 Processing System** and add it to the design.

2.1.5 Click *Run Block Automation*, and click **OK** to accept the default settings.

2.1.6 Double click on the Zynq block to open the *Customization* window for the Zynq processing system.

A block diagram of the Zynq should now be open, showing various configurable blocks of the Processing System.

### 2.2 Configure the I/O Peripherals block to only have UART 1 support.

2.2.1 Click on the *MIO Configuration* panel to open its configuration form.

2.2.2 Expand the *I/O Peripherals* on the right.

2.2.3 Uncheck *ENET 0*, *USB 0*, and *SD 0*, *GPIO (GPIO MIO)*, leaving *UART 1* selected.

2.2.4 In the **MIO Configuration** panel, expand the **Application Processing Unit** and uncheck the **Timer 0**.

2.2.5 Click **OK**.

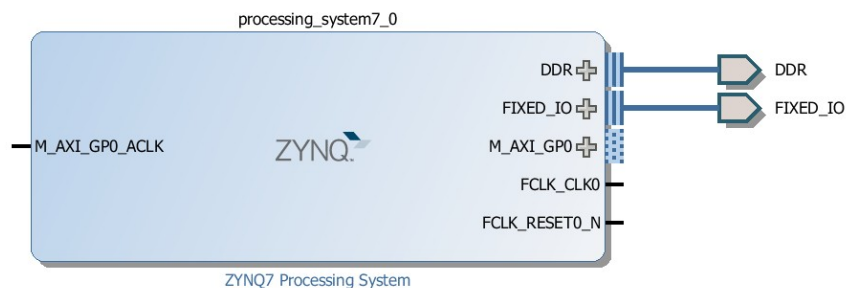




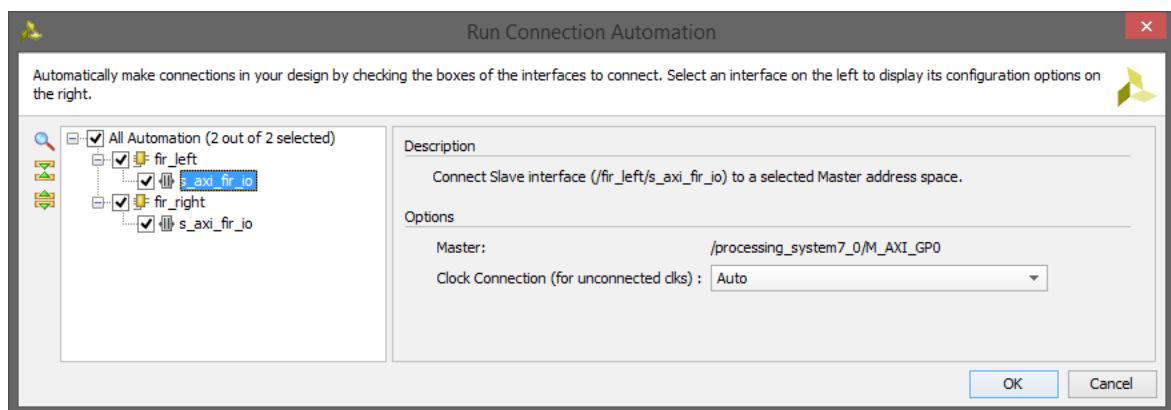
Figure 2. ZYNQ Processing System configured block

## 3Add FIR Core to the System

## Step 3

### 3.1 Instantiate the provided FIR core twice naming the instances as **fir\_left** and **fir\_right**. Validate the design.

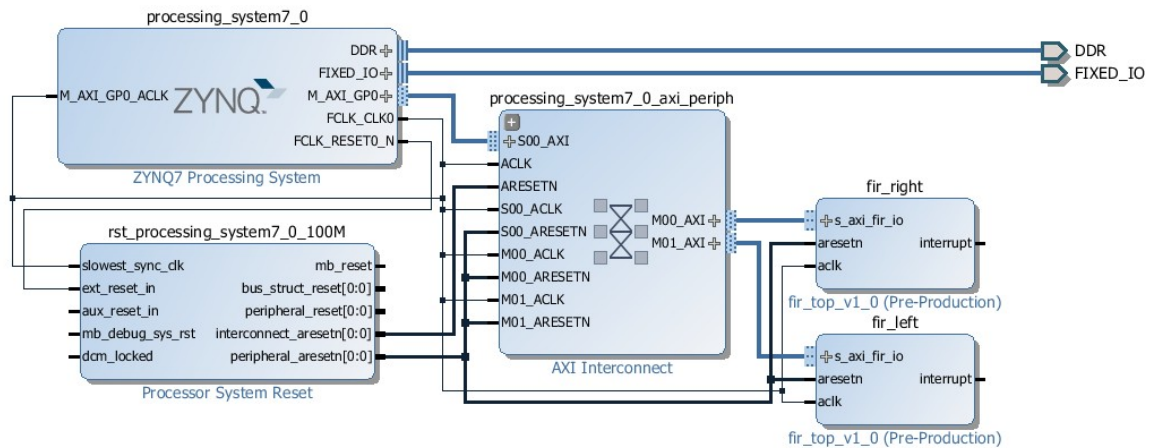
- 3.1.1 Click the Add IP icon  and search for **fir** in the catalog.
- 3.1.2 Double-click on the **fir\_top\_v1\_0** to add the IP instance to the system
- 3.1.3 Select the *fir\_top\_1* instance and change its name to **fir\_left** in its property form.
- 3.1.4 Click the Add IP icon  and search for **fir** in the catalog.
- 3.1.5 Double-click on the **fir\_top\_v1\_0** to add the IP instance to the system
- 3.1.6 Select the *fir\_top\_0* instance and change its name to **fir\_right** in its property form.
- 3.1.7 Click on **Run Connection Automation**, and select **All Automation** to select *fir\_left* and *fir\_right*.
- 3.1.8 Click on *s\_axi\_fir\_io* for both *fir\_left* and *fir\_right* and confirm that they will be automatically connected to the Zynq *M\_AXI\_GP0* port



**Figure 3. Connection automation for the FIR IP blocks**

- 3.1.9 Click **OK** to connect the two blocks to the *M\_AXI\_GP0* interface.

The design should look similar to shown below:



**Figure 4. The completed design**

It is not necessary to connect the *interrupt* signals of the *fir* blocks.

- 3.1.10** Select **Tools > Validate Design** to run the design rule checker and to make sure that there are no design errors.

## 4 Generate the Bitstream

### Step 4

### 4.1 Create the top-level HDL of the embedded system, and generate the bitstream..

- 4.1.1** In Vivado, select the *Sources* tab, expand the *Design Sources*, right-click the *system.bd* and select **Create HDL Wrapper** and click **OK**.
- 4.1.2** Click on the **Generate Bitstream** in the *Flow Navigator* pane to synthesize and implement the design, and generate the bitstream.
- 4.1.3** Click **Save** to save the design and **Yes** to run the necessary processes.
- 4.1.4** During the bitstream generation process click “Open Implemented Design” if asked.

## 5 Export the Design to the Vitis

### Step 5

### 5.1 Export the design to the Vitis, create the software XSA using the standalone operating system and enable the profiling options.

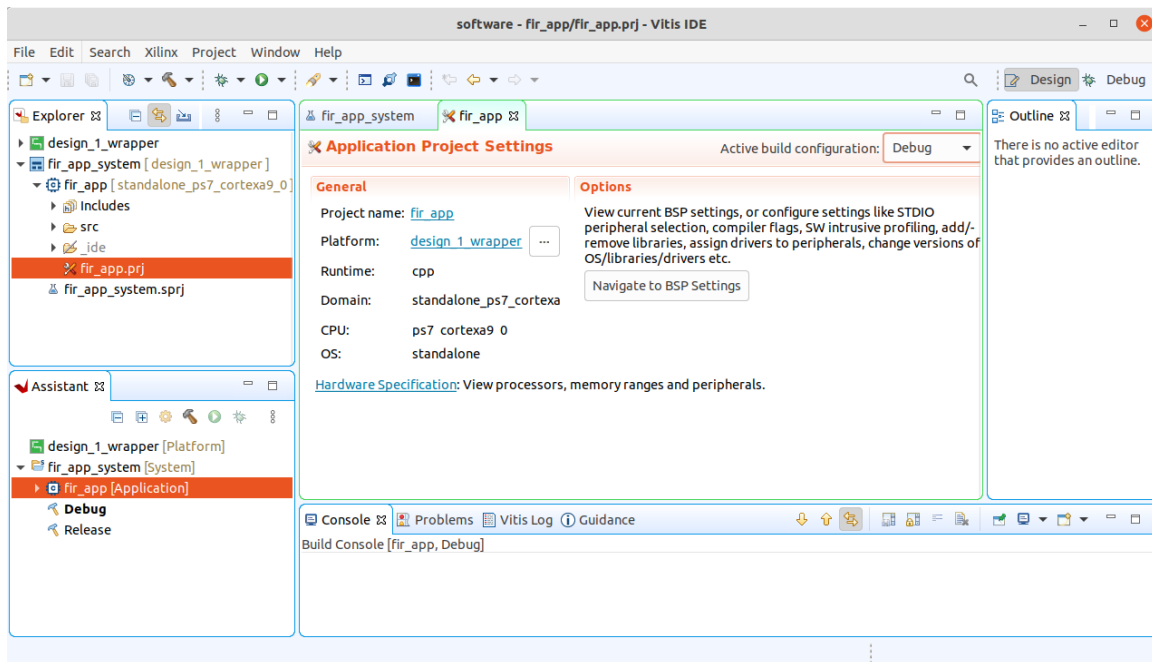
- 5.1.1** Export the hardware configuration by clicking **File > Export > Export Hardware...**
- 5.1.2** Tick the box to *Include Bitstream*, and click **OK**
- 5.1.3** Launch Vitis by clicking **File > Launch Vitis IDE** and click **OK**

5.1.4 In Vitis, select **File > New > Application Project**.

5.1.5 Select you exported XSA file from the hardware

5.1.6 Notice **Standalone\_bsp\_0** in the **Project name** field and click **Finish** with default settings.

A Board Support Package Settings window will appear.



- 5.1.7 Click on **Navigate to BSP setting**. Then click **modify BSP setting**. Select the **Overview > standalone** entry in the left pane, click on the drop-down arrow of the *enable\_sw\_intrusive\_profiling* Value field and select **true**.

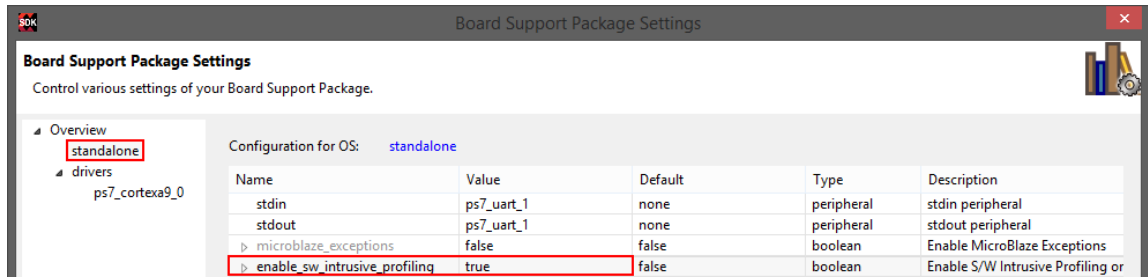


Figure 5. Enable profiling in the board support package

- 5.1.8 Select the **Overview > drivers > ps7\_cortexa9\_0** and add **-pg** in addition to the **-g** in the *extra\_compiler\_flags* Value field.

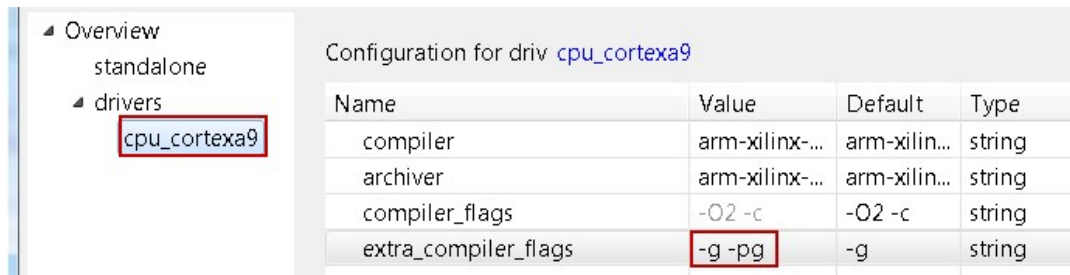


Figure 6 Adding profiling switch

- 5.1.9 Click **OK** to accept the settings and create the BSP.

## 6 Create the Application

## Step 6

- 6.1 Create the *lab* application using the provided *lab7.c*, *fir.c*, *fir.h*, *fir\_coef.dat*, and *xfir\_fir\_io.h* files.

- 6.1.1 Select **File > New > Application Project**.

- 6.1.2 Enter a name for the project name, select the **Use existing standalone\_bsp\_0** option, and click **Next**.

- 6.1.3 Select **Empty Application** in the *Available Templates* pane and click **Finish**.

- 6.1.4 In the *lab7* project, right click on the *src* directory and select **Import**.

- 6.1.5 Expand the **General** folder and double-click on **File system**, and browse to the *lab\_sources* directory.

- 6.1.6 Select *fir\_coef.dat*, *fir.c*, *fir.h*, *lab.c*, *lab.h* and *xfir\_fir\_io.h*, and click **Finish**.

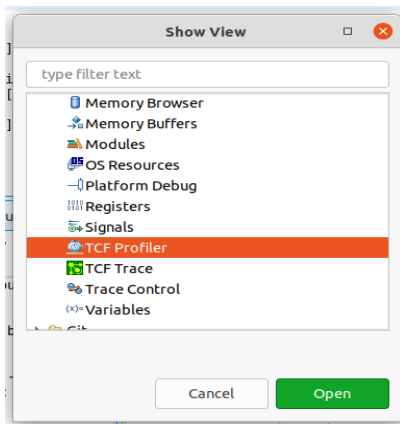
The program should compile successfully and generate the .elf file.

**6.1.7** Open the *lab.c* file and scroll to the main function at the bottom. Notice the following code:

```
#ifdef SW_PROFILE
    fir_software(&output,signal);
#else
    filter_hw_accel_input(&output,signal);
#endif
```

The function *fir\_software()* is a software implementation of the FIR function. The *filter\_hw\_accel\_input()* function offloads the FIR function to the two FIR blocks that have been implemented in the PL. In this case, you are using the implementation of the FIR filter on hardware (PL) as this gray code disabled.

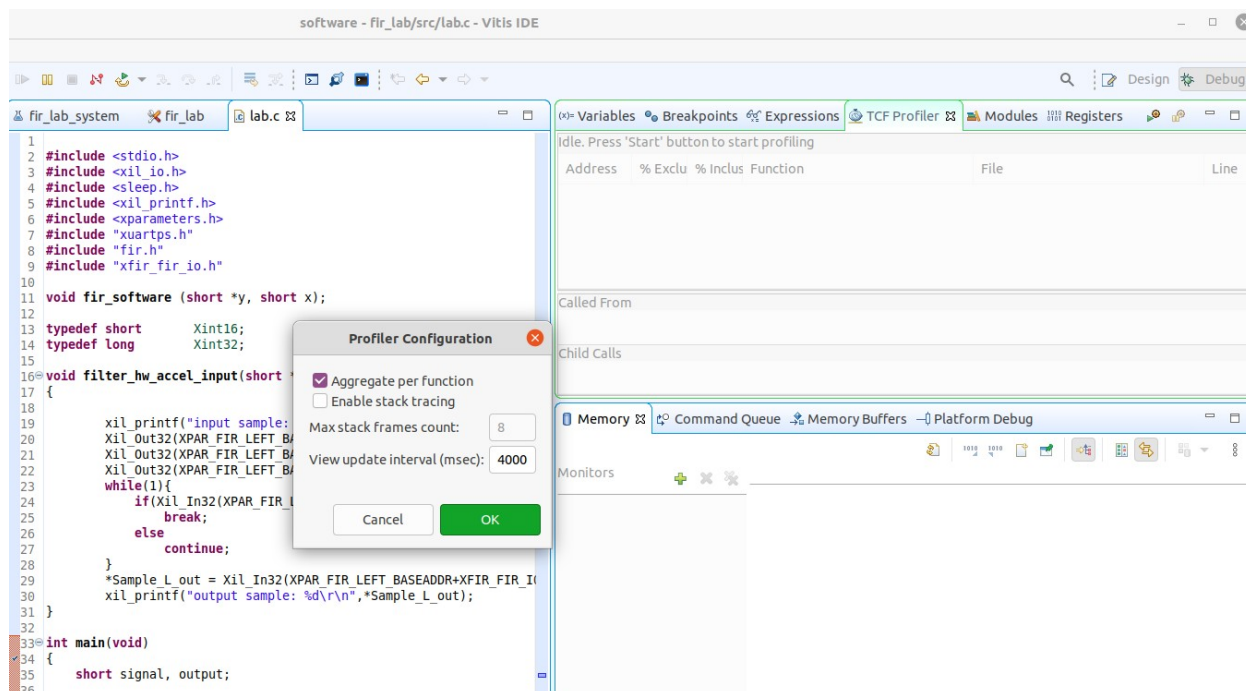
**6.1.8 Run TCP Profiler.** Select Window > Show View > Debug > TCP Profiler



**6.1.9** Start the application

**6.1.10** Start the TCF profiler





### 6.1.11 Observe the output of the TCF profiler

Address	% Exclu	% Includ	Function	File	Line
00100f0c	.000	100	_start	xil-crt0.S	59
00101414	.000	100	__libc_fini_array	fini.c	31
00101c48	100	100	_exit	_exit.c	16

## 7 Run the Application and Profile

## Step 7


**7.1 Place the board into the JTAG boot up mode. Program the PL section and run the application using the user defined SW\_PROFILE symbol.**

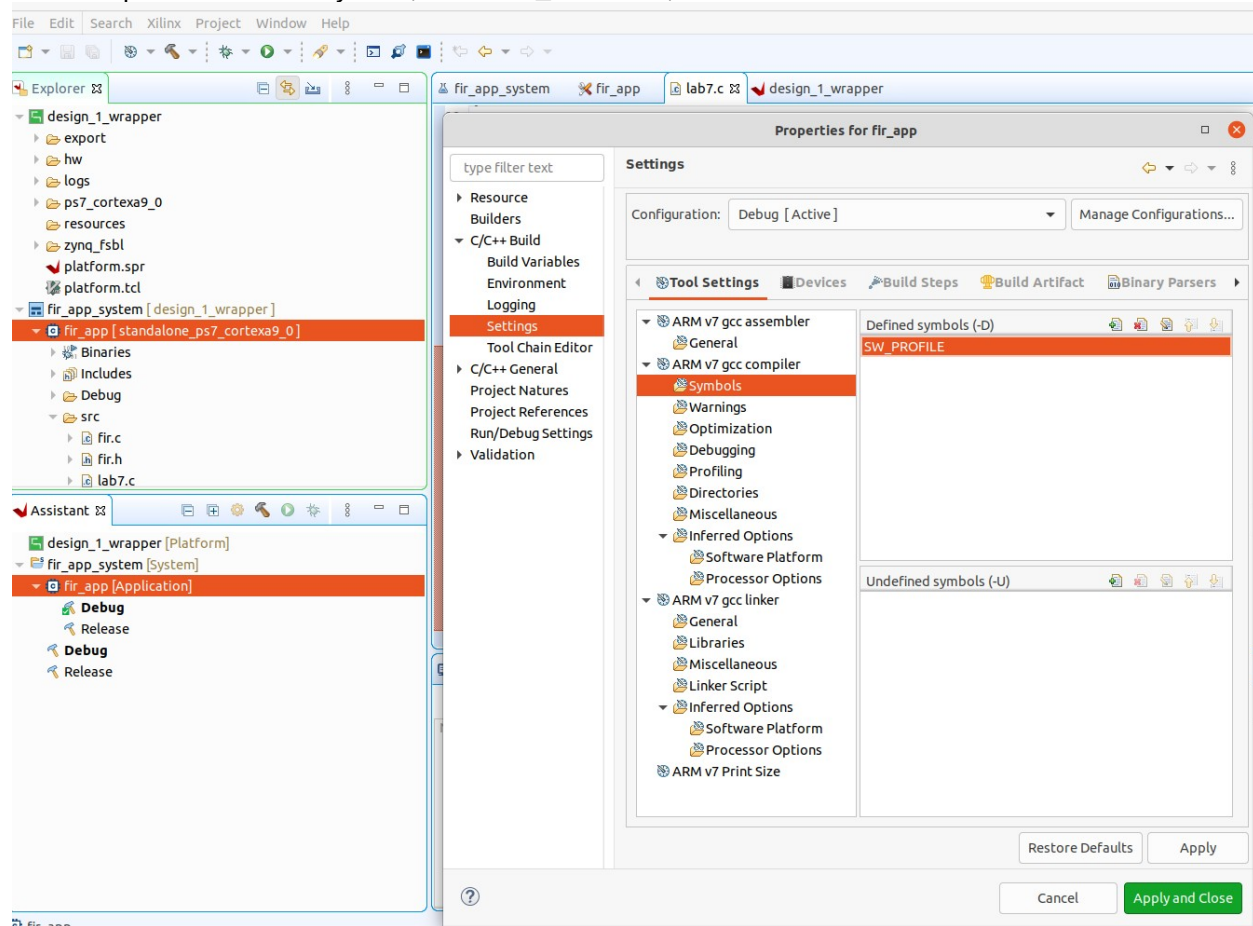
**7.1.1** Place the board in the JTAG boot up mode.

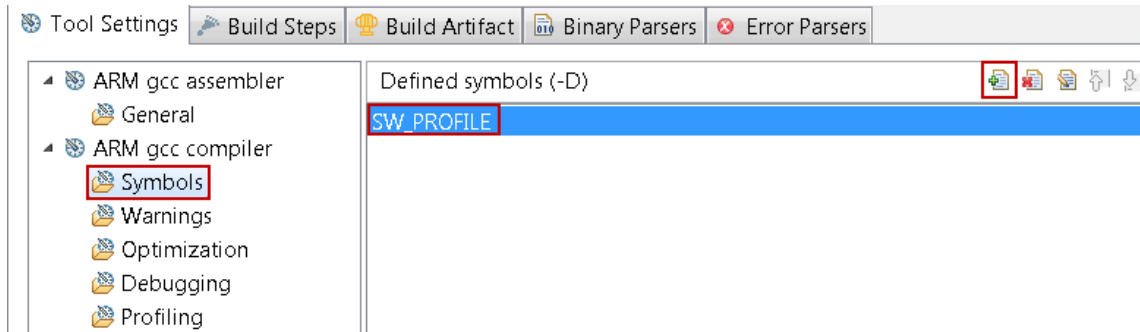
**7.1.2** Power ON the board.

**7.1.3** Select **Xilinx Tools > Program FPGA** and click on **Program**.

**7.1.4** Right click on the *lab directory*, and select **C/C++ Build Settings**.

**7.1.5** Under the **ARM gcc compiler** group, select the **Symbols** sub-group, click on the  button to open the value entry form, enter **SW\_PROFILE**, and click **OK**.



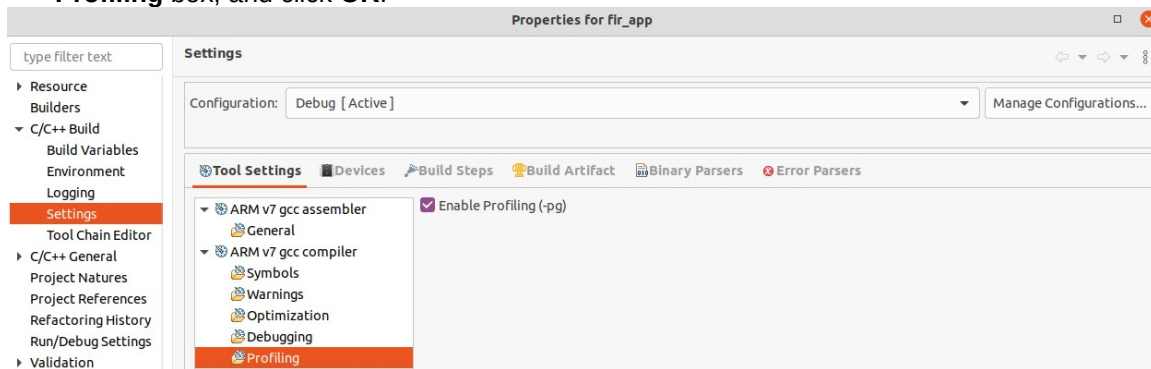


This will allow us to profile the software loop of the FIR application. You can verify this by looking at lab7.c source code to verify that the proper part of the code is inactive.

Q2. What change do you see in lab7.c source file as a result of this step?

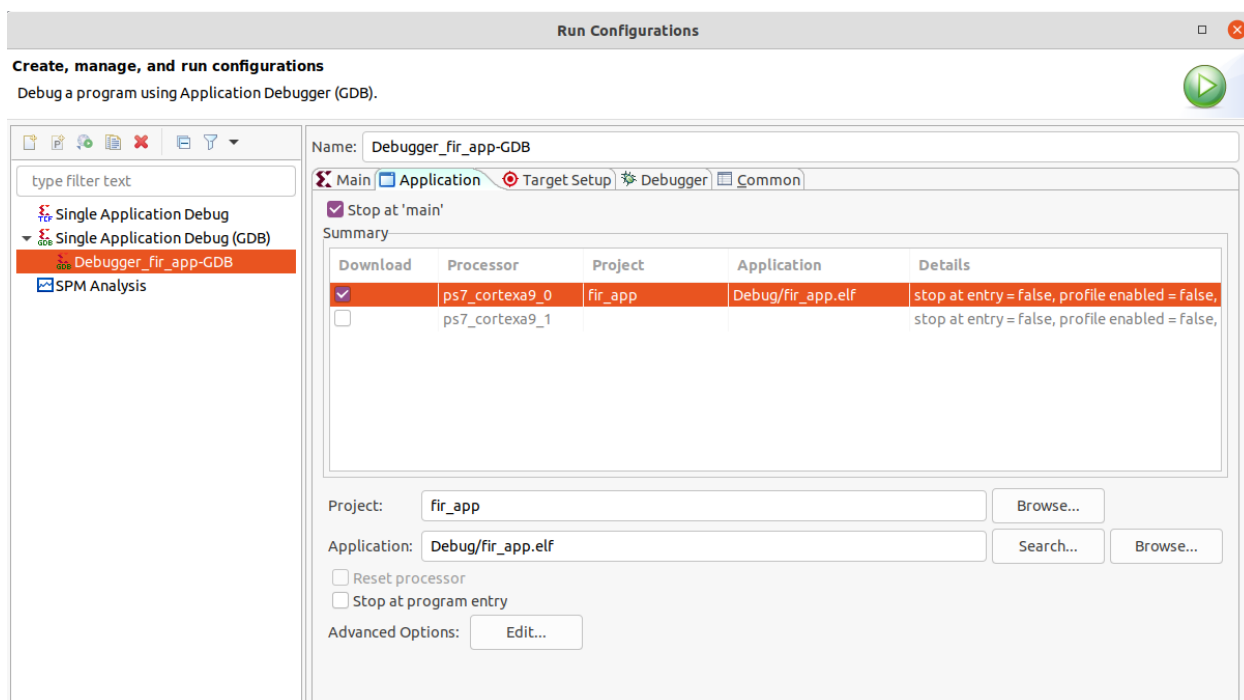
**Figure 7. Add user-defined symbol**

- 7.1.6** Under the **ARM gcc compiler** group, select the **Profiling** sub-group, then check the **Enable Profiling** box, and click **OK**.

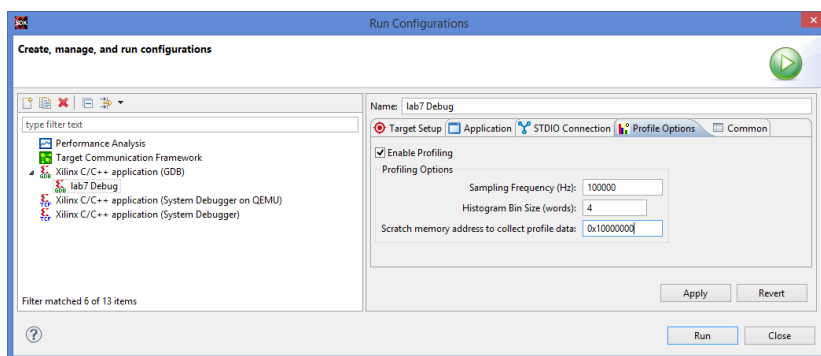


**Figure 8. Compiler setting for enabling profiling**

- 7.1.7** From the menu bar, Select **Run > Run Configurations...** and double click on *Xilinx C/C++ application* to create a new configuration.
- 7.1.8** Click on the newly created **lab Debug** configuration, and select the **Application** tab and press **Advanced Options: Edit**



- 7.1.9** Click on the *Enable Profiling* check box, enter **100000** (100 kHz) in the Sampling Frequency field, enter **0x10000000** in the scratch memory address field as shown in Figure 9, and click **Apply**.



**Figure 9. Profiling options**

- 7.1.10** Click the **Run** button to download the application and execute it.

The program will run, and when execution has completed, a message will be displayed indicating that the profiling results are being saved in gmon.out file at the lab\Debug directory.

- 7.1.11** Click **OK**.

## 7.2 Invoke gprof and analyze the results.

- 7.2.1** Expand the *Debug* folder under the **lab7** project in the Project Explorer view, and double click on the **gmon.out** entry.

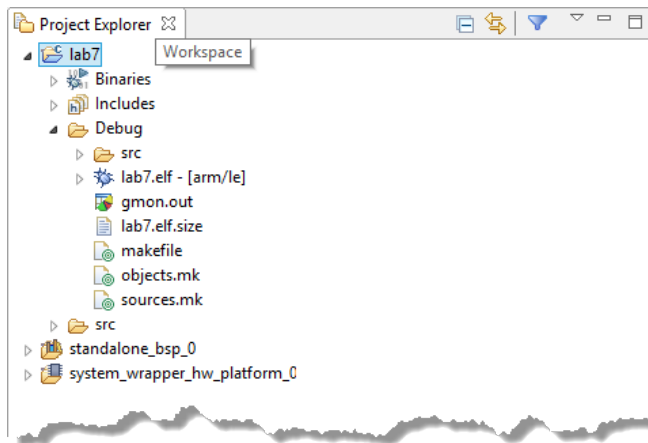


Figure 10. Invoking gprof on gmon.out

7.2.2 The Gmon File Viewer dialog box will appear showing *lab7.elf* as the corresponding binary file. Click **OK**. This will open gprof window on the bottom right corner of the screen.

7.2.3 Click on the **Sort samples per function** button (  ).

7.2.4 Click in the **%Time** column to sort in the descending order.

Note that the *fir\_software* routine is called 60 times, 21 samples were taken during the profiling, and on an average of 3.499 microseconds were spent per call.

Name (location)	Samples	Calls	Time/Call	% Time
Summary	23			100.0%
fir_software	21	60	3.499us	91.3%
mcount	1			4.35%
memcpy	1			4.35%
cortexa9_init	0	0		0.0%
main	0	0		0.0%
XScuGic_DeviceInitialize	0	1	0ns	0.0%
XScuGic_RegisterHandler	0	1	0ns	0.0%

Figure 11. Sorting results

7.2.5 Go back to the *Run Configuration*, and change the sampling frequency to **10000** (10 KHz), **1000000** (1 MHz), and **5000000** (5 MHz) and profile the application again.

7.2.6 For each case invoke **gprof**, select the **Sorts samples per function** output, and sort the **%Time** column.

Q3. Present a table/graph and compare your results. Write your conclusion from applying different sampling frequency to the profiling process. Is there any relationship between the number of samples taken and number of function calls? How about other factors? Explain exactly how profiling is affected by sampling frequency? What is more accurate; lower sampling rate or higher sampling rate? What happens if sampling rate is decreases to its minimum possible? What happens if sampling rate is decreases to its maximum possible?

At this stage, the designer of the system would decide if the FIR function should be ported to hardware.

### 7.3 Profile the application using the hardware FIR filter IP by removing the user defined `SW_PROFILE` symbol.

**7.3.1** Select the *lab7* application, right-click, and select **C/C++ Build Settings**.

**7.3.2** Under the **ARM gcc compiler** group, select the **Symbols** sub-group, select `SW_PROFILE`, and delete it by clicking on the delete button.

This will allow us to profile the hardware IP of the FIR application. You can verify this by looking at *lab7.c* source code to verify that the proper part of the code is inactive.

Q4. What change do you see in *lab7.c* source file as a result of this step?

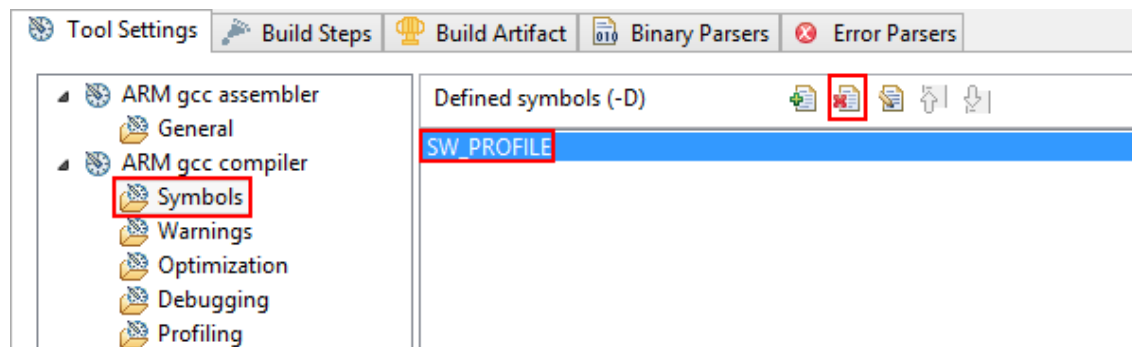


Figure 12. Deleting the user-defined symbol

**7.3.3** Click **Apply**, and then click **OK**

**7.3.4** Select **Run > Run Configurations** and change the sampling frequency to **10000** (10 KHz), **100000** (100 KHz), **1000000** (1 MHz), and **5000000** (5 MHz) and profile the application again.

**7.3.5** Open *gmon.out*, select the **Sorts samples per function** output, and sort the %Time column.

Q5. Write your conclusion from applying different sampling frequency to the profiling process. Is there any relationship between the number of samples taken and number of function calls? How about other factors? Explain exactly how profiling is affected by sampling frequency? What is more accurate; lower sampling rate or higher sampling rate? What happens if sampling rate is decreases to its minimum possible? What happens if sampling rate is decreases to its maximum possible?

Q6. How does software profiling compare to hardware profiling? Elaborate your answer. Has the number of calls changed? Why or why not? Has the time changed? Why or why not?

Q6. What is the most noticeable change in the hardware profile compared to software profile?

Notice that the output now shows *filter\_hw\_accel\_input* function call instead of the *fir\_software* function call. Note that the number of calls to the filter function has not changed but the average time spent per call is 1.799 us as the filtering is done in the hardware instead of the software.

Also notice that the amount of time spent in the filtering function reduced from about 98% to 0.26%.

Q7. What is the reason for this significant reduction in average time spent per call as well as time spent for filtering function? Explain your answer. Find the ratio of the speed up for each sampling frequency while profiling the application. Explain your result using a table or graph.

Name (location)	Samples	Calls	Time/Call	% Time
Summary	42184			100.0%
> XUartPs_SendByte	40181	2248	17.874us	95.25%
> Xil_In32	696			1.65%
> Xil_In16	348			0.82%
> Xil_Out8	348			0.82%
> mcount	189			0.45%
> XUartPs_RecvByte	124			0.29%
> filter_hw_accel_input	108	60	1.799us	0.26%
> outnum	54	0		0.13%

**Figure 13. Profiling the application with the hardware IP**

**7.3.6** Close the SDK and Vivado programs by selecting **File > Exit** in each program.

**7.3.7** Turn OFF the power on the board.

Q8. Explain how FIR filter is integrated to the SDK project as a hardware accelerated HLS core? How does software communicate with the FIR core on FPGA fabric? What is being sent to the FIR core as input data and what do you expect to see as output?

Q9. Write down your conclusion from this lab.